

Programming Assignment # 1

Due: Feb 27

In-Order RV-32IF (Integer and Floating Point) Pipeline

In this assignment, you will become familiar with RISC-V integer and floating point pipelines. You will create a cycle-accurate simulator of an in-order RISC-V processor, called `rvsim`, that supports the RV-32IF ISA—*i.e.*, 32-bit RISC-V with integer and single-precision floating point instructions. Your simulator will be very basic: it will model all the necessary pipeline interlocks to enforce correct pipeline execution, but it will not support out-of-order execution, data forwarding, nor branch prediction.

1 Files

The first thing you should do is download the files from ELMs. There are 18 files: `asm.c`, `main.c`, `pipeline.c`, `pipeline.h`, `fu.c`, `fu.h`, `output.c`, `output.h`, `Makefile`, `io_pipe.fu`, `simple.s`, `simple.io.out`, `vect.s`, `vect.io.out`, `newton.s`, `newton.io.out`, `cos.s`, and `cos.io.out`.

`asm.c` is an assembler for the RV-32IF ISA which your simulator will implement (more about the assembler and the ISA later). `main.c`, `pipeline.c`, `pipeline.h`, `fu.c`, `fu.h`, `output.c`, and `output.h` are the simulator source files. Much of the simulator's internals, like the functional units and the simulator's output generator, have been implemented already and are provided in these files. Your assignment is to build the main control module of the pipeline (most of the changes will go in `pipeline.c`). `Makefile` is a unix make file which will produce the binaries `asm` and `rvsim`. And `io_pipe.fu` is a functional unit configuration file (more about this later).

In the remaining files, we have provided four example assembly programs. `simple.s` performs a simple arithmetic sequence, `vect.s` performs a vector operation, `newton.s` finds the left-most root of a quadratic expression using Newton's method of root approximation, and `cos.s` estimates the cosine of "1" using a Taylor series expansion with 10 iterations. Finally, `simple.io.out`, `vect.io.out`, `newton.io.out`, and `cos.io.out` contain the output for the 4 benchmarks, respectively, on a properly functioning simulator.

2 RV-32IF ISA

You will be simulating the RV-32IF ISA from Hennessy & Patterson. This is a 32-bit ISA: all registers and data paths are 32 bits wide, and all instructions will operate on 32-bit operands. Within this 32-bit data size, you will support both signed and unsigned fixed point data types as well as single-precision floating point data types. To simplify the project, you will not implement the RV-32IF spec fully. In particular, you will only support a subset of the RV-32IF instructions.

Integer Arithmetic / Logic			
add rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] + \text{reg}[\text{rs2}]$	addi rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] + \text{imm}$
sub rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] - \text{reg}[\text{rs2}]$		
sll rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \ll \text{reg}[\text{rs2}]$	slli rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \ll \text{imm}$
srl rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \gg \text{reg}[\text{rs2}]$	srli rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \gg \text{imm}$
and rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \& \text{reg}[\text{rs2}]$	andi rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \& \text{imm}$
or rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \text{reg}[\text{rs2}]$	ori rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \text{imm}$
xor rd rs1 rs2	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \wedge \text{reg}[\text{rs2}]$	xori rd rs1 imm	$\text{reg}[\text{rd}] = \text{reg}[\text{rs1}] \wedge \text{imm}$
slt rd rs1 rs2	$\text{reg}[\text{rd}] = (\text{reg}[\text{rs1}] < \text{reg}[\text{rs2}])$	slti rd rs1 imm	$\text{reg}[\text{rd}] = (\text{reg}[\text{rs1}] < \text{imm})$
sltu rd rs1 rs2	$\text{reg}[\text{rd}] = (\text{reg}[\text{rs1}] < \text{reg}[\text{rs2}])$	sltiu rd rs1 imm	$\text{reg}[\text{rd}] = (\text{reg}[\text{rs1}] < \text{imm})$
Floating Point Arithmetic		Memory	
fadd.s rd rs1 rs2	$\text{fpreg}[\text{rd}] = \text{fpreg}[\text{rs1}] + \text{fpreg}[\text{rs2}]$	lw rd rs1 imm	$\text{reg}[\text{rd}] = \text{Mem}[\text{reg}[\text{rs1}] + \text{imm}]$
fsub.s rd rs1 rs2	$\text{fpreg}[\text{rd}] = \text{fpreg}[\text{rs1}] - \text{fpreg}[\text{rs2}]$	sw rs2 rs1 imm	$\text{Mem}[\text{reg}[\text{rs1}] + \text{imm}] = \text{reg}[\text{rs2}]$
fmult.s rd rs1 rs2	$\text{fpreg}[\text{rd}] = \text{fpreg}[\text{rs1}] * \text{fpreg}[\text{rs2}]$	flw rd rs1 imm	$\text{fpreg}[\text{rd}] = \text{Mem}[\text{reg}[\text{rs1}] + \text{imm}]$
fdiv.s rd rs1 rs2	$\text{fpreg}[\text{rd}] = \text{fpreg}[\text{rs1}] / \text{fpreg}[\text{rs2}]$	fsw rs2 rs1 imm	$\text{Mem}[\text{reg}[\text{rs1}] + \text{imm}] = \text{fpreg}[\text{rs2}]$
Control			
jal rd off	$\text{reg}[\text{rd}] = \text{pc} + 4, \text{pc} = \text{pc} + \text{off}$		
jalr rd r1	$\text{reg}[\text{rd}] = \text{pc} + 4, \text{pc} = \text{reg}[\text{r1}]$		
beq rs1 rs2 imm	$\text{pc} = (\text{reg}[\text{rs1}] == \text{reg}[\text{rs2}]) ? \text{pc} + \text{imm} : \text{pc} + 4$		
bne rs1 rs2 imm	$\text{pc} = (\text{reg}[\text{rs1}] != \text{reg}[\text{rs2}]) ? \text{pc} + \text{imm} : \text{pc} + 4$		
halt	stop the simulation		

Table 1: Instructions from the RV-32IF ISA supported by the simulator.

Table 1 lists the instructions you will implement, including the assembly notation and the C-expression for each instruction. Also, although immediate values in RV-32I branch and jump instructions are left-shifted by 1 (see Hennessy & Patterson), your control instructions **SHOULD NOT** perform the left shift.

The datatype operated upon by each instruction depends on the opcode. The integer arithmetic instructions `add`, `sub`, and `addi` operate on signed integer values. The logic instructions `sll`, `slli`, `srl`, `srli`, `and`, `andi`, `or`, `ori`, `xor`, and `xori` operate on unsigned integer values. And, the floating point arithmetic instructions `fadd.s`, `fsub.s`, `fmult.s`, and `fdiv.s` operate on floating point values. The comparison instructions have both signed and unsigned integer versions: `slt` and `slti` operate on signed integer values whereas `sltu` and `sltiu` operate on unsigned integer values. (Unlike other ISAs, RISC-V does not provide unsigned arithmetic integer instructions—*i.e.*, there is no `addu` instruction. Instead, unsigned arithmetic is performed using signed operators, and then overflow is detected by using the unsigned comparison operators.)

Notice all the instructions in Table 1 except for one exist in the normal RV-32IF ISA. The instruction we’ve added is `halt`. As its name implies, when your simulator executes a `halt` instruction, it should terminate the simulation.

3 asm: An Assembler for the RV-32IF ISA

We have provided an assembler, `asm.c`, that supports all the instructions in Table 1. The `asm.c` file is fully functional, and you will not need to make any modifications to this file. Simply use the `Makefile` to make the binary `asm` from the `asm.c` source file.

The format for assembly programs is very simple. A valid assembly program is an ASCII file in which each line of the file represents a single instruction, or a data constant. The format for a single line of assembly code is:

```
label<tab>instruction<tab>field0<tab>field1<tab>field2<tab>comments
```

Here's a simple example of a full assembly program:

```

start    addi    x1    x0    #5    load reg[1] with 5
         addi    x1    x1    #-1   decrement reg[1]
         flw     f0    x0    var1   loads fpreg[0] with value stored in var1
         fadd.s  f2    f0    f0     double fpreg[0]
         fsw     f2    x0    var1   put fpreg[0] back
         beq     x1    x0    done   goto done when reg[1]==0
         jal     x0    start
done     halt
var1     .df     32.0               Declare a variable, initialized to 32

```

The leftmost field on a line is the label field which indicates a symbolic address. Valid labels contain a maximum of 6 characters and can consist of letters and numbers. The label is optional (the tab following the label field is not). After the optional label is a tab. Then follows the instruction field, where the instruction can be any one of the assembly-language mnemonics described in Section 2. After another tab comes three fields, also separated by tabs. Each field represents either a register identifier or an immediate value. Register identifiers begin with the letter “x” or “f” as in “x0” or “f0” for integer and floating point registers, respectively. Immediates are specified as a constant decimal value preceded by a “#” sign as in “#5.” For the memory and control instructions, the `imm` field can either be a decimal value preceded by the “#” sign, or a label can be used as in “var1.” After the last field is another tab, then any comments. The comments end at the end of the line.

In addition to instructions, lines of assembly code can also include directives for the assembler. A directive tells the assembler to put a constant value into the place where an instruction would normally be stored. We support two directives, “.dw” and “.df,” for storing integer and floating point constants, respectively.

4 Pipeline Simulator

Your assignment is to create a cycle-accurate simulator of an in-order RV-32IF processor. We have provided some code to get you started in `main.c`, `fu.c`, and `output.c`. This code will make your life easier since you won't have to write the whole simulator from scratch, and it will enforce some coding disciplines that ensure you actually simulate the details of the processor pipeline. The following sections describe both the code we provide and the code you will write in greater detail.

4.1 Pipeline Structure

We have provided pipeline register data structures specified in the file, `pipeline.h`. To enforce proper pipeline simulation, you should use these structures in your simulator. Here is the main processor structure, `state_t`:

```
typedef struct _state_t {
    /* memory */
    unsigned char mem[MAXMEMORY];

    /* register files */
    rf_int_t rf_int;
    rf_fp_t rf_fp;

    /* pipeline registers */
    unsigned long pc;
    if_id_t if_id;

    fu_int_t *fu_int_list;
    fu_fp_t *fu_add_list;
    fu_fp_t *fu_mult_list;
    fu_fp_t *fu_div_list;

    wb_t int_wb;
    wb_t fp_wb;

    int fetch_lock;
} state_t;
```

Both the integer and floating point pipelines pass through four types of pipeline stages: fetch, decode, execute, and writeback. The “pc” field and “if_id” structure contain the pipeline registers for the fetch and decode stages, respectively. These pipeline registers are shared by both integer and floating point pipelines, so only one instruction (integer or floating point) can be fetched and decoded per cycle. The “fu_int_list” pointer points to the pipeline registers for the integer execute stage(s), and the “fu_add/mult/div_list” pointers point to the pipeline registers for the floating point add, multiply, and divide execute stages. (We will discuss these pipeline registers further in Section 4.2). And finally, the “int_wb” and “fp_wb” structures contain the pipeline registers for the writeback stages of the integer and floating point pipelines, respectively. At most 1 integer and 1 floating point instruction (*i.e.*, 2 instructions total) can writeback per cycle.

In addition to the pipeline registers, the processor state also contains memory, “mem,” and two register files—one for the integer pipeline, “rf_int,” and another for the floating point pipeline, “rf_fp.” Each register file contains an array of registers. Notice the integer registers are of type “int_t,” as defined in `fu.h`, which allows for a signed or unsigned interpretation of each register’s contents through a “union” construct. This permits both signed and unsigned operations on the integer register data.

4.2 Simulator Loop, Functional Units, and Output Generator

Included in the download files are three modules that contain code we provide: `main.c`, `fu.c`, and `output.c`. This section describes the functionality implemented in these modules. Note, while the code we provide is fully functional, you are welcome to modify most of it. In fact, to complete the assignment, you will need to make at least a few modifications to our code in order to integrate it with your code. The only code you *cannot* modify is the output generator code which we will discuss in Section 4.2.3.

4.2.1 `main.c`

`main.c` contains the main simulator loop. This loop calls the `fetch`, `decode`, `execute`, and `writeback` functions you will provide to simulate each type of pipeline stage (see Section 4.3). One iteration through the main simulator loop advances each pipeline stage by 1 cycle, and corresponds to 1 simulated processor cycle. Notice the pipeline stages are called in reverse order. This ensures each instruction propagates only 1 pipeline stage per cycle. (Calling the stages in forward order would allow a new instruction to propagate all the way down the pipeline in 1 iteration). `main.c` also contains a function to parse the command line arguments. The simulator expects two arguments: the assembled program file preceeded by the flag “-b,” and the functional unit configuration file preceeded by the flag “-o.”

4.2.2 `fu.c`

`fu.c` contains a significant amount of code that implements the integer/floating point functional units and the simulation of their execute stages. The functional unit implementation consists of 4 types of routines. First, `state_create` reads a functional unit configuration (FU config) file, and creates and initializes the functional units accordingly. We have provided an FU config file, called `io_pipe.fu`. Each line in the FU config file specifies a functional unit to be simulated, providing the name and type of functional unit. There are four valid functional unit types: INT, ADD, MULT, and DIV (each type has a corresponding linked list in the pipeline register data structure described in Section 4.1). ADD executes the `fadd.s` and `fsub.s` instructions, MULT executes the `fmult.s` instruction, and DIV executes the `fdiv.s` instruction. All remaining instructions are of type integer and execute in INT-type functional units. Each line in the FU config file also specifies a list of numbers, one for each pipeline stage that specifies the number of cycles in that stage. For example, the floating point multiplier in `io_pipe.fu` contains 4 1-cycle stages.

Second, an instruction can be “issued” into one of the functional units via the `issue_fu_int` or `issue_fu_fp` routines. The former is used for integer instructions, and the latter for floating point instructions. These routines return “0” if the instruction is successfully issued; otherwise, they return “-1” indicating no functional units of that type are free (*i.e.*, a structural hazard). Third, the functional units can be advanced one cycle by calling the `advance_fu_int` and `advance_fu_fp` routines. These routines move previously issued instructions through the functional units’ execute stages. When execution of an instruction completes, the instruction is placed into the appropriate writeback stage pipeline register, `int_wb` or `fp_wb`, depending on the instruction’s type (integer or floating point, respectively). Finally, the `fu_int_done` and `fu_fp_done` routines test for instructions still in-flight in functional units of type integer or floating point, respectively, and return TRUE when no instructions are in flight or FALSE when one or more instructions are still in flight.

In addition to implementing the functional units, `fu.c` also provides a routine for decoding instructions, `decode_instr`, which takes two arguments, an instruction and a pointer to a flag (`use_imm`), and returns a pointer into the instruction table defined at the top of `fu.c`. Each instruction table entry provides 4 pieces of information: name, group number, operation number, and data type. The group number specifies a group of instructions from Table 1 to which the instruction belongs. `FU_GROUP_INT` contains the instructions in the “Integer Arithmetic / Logic” portion of Table 1. Within this group, the `use_imm` flag specifies whether the instruction uses an immediate value (*i.e.*, the right half of the “Integer Arithmetic / Logic” portion of Table 1), or whether the instruction does not use an immediate value (*i.e.*, the left half of the “Integer Arithmetic / Logic” portion). `FU_GROUP_ADD`, `FU_GROUP_MULT`, and `FU_GROUP_DIV` contain the instructions in the “Floating Point Arithmetic” portion of Table 1 (note, both `fadd.s` and `fsub.s` belong to `FU_GROUP_ADD`). `FU_GROUP_MEM` contains the instructions in the “Memory” portion of Table 1, and `FU_GROUP_BRANCH` contains the instructions in the “Control” portion of Table 1, excluding the `halt` instruction which has its own group, `FU_GROUP_HALT`. Within a group, the operation number specifies the actual decoded instruction. Finally, the data type specifies either integer (`DATA_TYPE_W`) or floating point (`DATA_TYPE_F`). We have provided the routine, `perform_operation`, which shows an example use of `decode_instr`. In `perform_operation`, the `add` instruction is fully decoded and implemented.

Notice `decode_instr` uses three macros, `FIELD_OPCODE`, `FIELD_FUNC3`, and `FIELD_FUNC7` to extract opcode fields from instructions. These macros are defined in `fu.h` along with others that facilitate extraction of register specifiers (`FIELD_RS1/RS2/RD`), extraction and sign extension of 12-bit immediates in I-type and S-type instructions (`FIELD_IMM_I` and `FIELD_IMM_S`, respectively) and 20-bit signed offsets (`FIELD_OFFSET`).

4.2.3 output.c

`output.c` contains the output generator routine, `print_state`, which is called once every iteration of the main simulation loop. This routine (along with other routines in `output.c`) dumps the state of memory, the register files, and the pipeline registers. Your final submitted simulator should use these routines in their unmodified form; otherwise, we will not be able to grade your simulator. Of course, you can make modifications to these routines during debugging, but be sure to remove your modifications before submitting your simulator.

4.3 Your Code

Your job is to write the routines that simulate the pipeline stages in `pipeline.c`: `fetch`, `decode`, `execute`, and `writeback`. You will also need to modify some of the code we provide to properly integrate your code.

`fetch` should go to memory and fetch the current instruction specified by the program counter, and place the fetched instruction into the `if_id` pipeline register. `decode` should examine the instruction in the `if_id` pipeline register, decode the instruction, and determine whether it can issue into the functional units. As long as the instruction has no data, control, or structural hazards (described below), it should issue; otherwise, it should stall. `execute` should advance all the functional units by one cycle. And `writeback` should write the result of any instruction in the

`int_wb` or `fp_wb` pipeline registers into the register file. The following describes in detail some of the issues your code must deal with.

4.3.1 Dealing with Data Hazards

In this simulator, you will *stall* all data hazards. (Although stalling leads to poor performance, the goal for this project is to implement simple mechanisms only). Data hazards should be detected and stalled in the decode stage.

There are two types of data hazards you need to detect and stall: RAW and WAW. (WAR hazards cannot occur since your pipeline does not perform late reads). A RAW hazard occurs when the instruction in decode reads a register written by an earlier instruction still executing in a functional unit. If a RAW hazard is detected, the instruction in decode must stall until the earlier instruction reaches its writeback stage. We will assume register writes happen on the first half of the clock cycle while register reads happen on the second half of the clock cycle, so when the earlier instruction reaches writeback, the dependent instruction in decode can issue on the same cycle. (In fact, since the pipeline is simulated in reverse order, writeback is performed before decode within the same simulated cycle).

A WAW hazard occurs when the instruction in decode writes a register written by an earlier instruction still executing in a functional unit, and the number of execute cycles remaining for the earlier instruction is larger than the latency of the instruction in decode. (If the instruction in decode were to issue immediately, the writes would reorder). If a WAW hazard is detected, the instruction in decode must stall until the number of execute cycles remaining for the earlier instruction is equal to the latency of the instruction in decode. Note, your simulator should stall WAW hazards for the minimum number of cycles only. (For example, you should not stall until the earlier instruction reaches its writeback stage). Furthermore, you can assume the integer pipeline is a fixed pipeline. Since WAW hazards can only occur in variable pipelines, WAW hazards will never occur through integer registers, only through floating point registers.

4.3.2 Dealing with Control and Structural Hazards

In this simulator, you will *stall* all control hazards. When a control instruction enters the decode stage and issues (after resolving any data and structural hazards), you should immediately set the `fetch_lock` flag in the `state_t` structure to `TRUE`. This prevents fetch and decode on subsequent cycles (see the simulator loop in `main.c`), though you should allow fetch of the sequential successor immediately following the control instruction. When the control instruction enters writeback, you should set `fetch_lock` back to `FALSE`. So, the processor stalls for the entire time the control instruction is in-flight. If the control instruction is not taken, you should not update the fetch stage. However, if the control instruction is taken, you must set the program counter to the target address and squash the instruction in `if_id` (*i.e.*, the sequential successor immediately following the control instruction that was incorrectly fetched).

In addition to data and control hazards, there are two types of structural hazards you must stall. Once again, detection and stall should occur in decode. First, you should stall if there is no functional unit available for the instruction in decode. This occurs whenever you try to issue an instruction, but `issue_fu_int` or `issue_fu_fp` return “-1.” Second, you must ensure that at most 1 floating point and 1 integer instruction writeback on the same cycle. A structural hazard of this

kind occurs when the number of execute cycles remaining in *any* earlier instruction still executing in a functional unit is one less than the latency of the instruction in decode, and both instructions write to the same register file. This kind of structural hazard can only occur in a variable pipeline (similar to the WAW hazard). So, structural hazards in the writeback stage need only be checked for the floating point register file.

4.3.3 Executing Instructions

The functional unit models described in Section 4.2.2 only simulate the timing of instructions; you must implement the function performed by each instruction separately. We have started this for you in `perform_operation`, but you must provide the majority of the code. One question is when should `perform_operation` occur? Since the code in `output.c` only prints the instruction field of each pipeline register, you can call `perform_operation` at any time. One possibility is to call it in the decode stage upon instruction issue, and modify the functional units to carry the result of the instruction all the way to writeback. Another possibility is to call `perform_operation` from within the functional units themselves, perhaps in the first execute stage. Yet another possibility is to carry the operands for the instruction through the functional units and call `perform_operation` during writeback. It's up to you.

4.3.4 NOP Instruction

There are many instructions that perform no operation (for example, any instruction that writes register `x0`). The RISC-V architecture picks one of these to be the “official” NOP instruction: `addi x0 x0 #0`. We have provided the encoding for this instruction in `fu.h`, defined as the constant `NOP`, and the code in `output.c` recognizes this instruction and prints “NOP” rather than the `addi` equivalent.

You should use the NOP instruction to initialize the pipeline register data structures, and to replace or “squash” the instruction in the `if_id` pipeline register following a taken branch. The nice thing about NOPs is that there is no need to execute them. Hence, you can throw away a NOP instruction once it is encountered in the decode stage. This saves having to execute and writeback NOPs, though you will need to fetch and decode them.

4.3.5 Halting

When your simulator encounters the `halt` instruction, it should end the simulation. Notice, however, any earlier instruction preceeding the `halt` should be allowed to complete. So, once you encounter a `halt` instruction in the decode stage, you can stop fetching and decoding; however, you should keep calling `execute` and `writeback` until all in-flight instructions drain from the pipeline and writeback. (Use the `fu_int_done` and `fu_fp_done` routines to detect when all instructions have drained from the functional units).

5 Grading

We will grade your programming assignment by comparing the output of your simulator on each of the four benchmarks against the output files `simple.io.out`, `vect.io.out`, `newton.io.out`, and

`cos.io.out`. So, your goal should be to design your simulator such that you get the exact same output as the output files we've provided.

5.1 Academic Honesty

Do not allow any other student to see any of your code. You may however discuss the assignment in general terms, with the other students. If copying or excessive collaboration is detected in your submissions, the matter will be referred to the Student Honor Counsel.